

Keywords: Moneyball, Baseball, Sabermetrics, Logistic Regression

Code supporting this analysis can be found at: <https://github.com/matthewwankiewicz/moneyball>

Introduction

In 2003, Michael Lewis released a book called Moneyball.¹ Moneyball is the story of the Oakland Athletics who, after losing players to the richer teams in the league, decided to focus on using the misfits of baseball to try to win a World Series. These “misfits” were players who never received support from teams because their traditional stats didn’t look good but they excelled in the stats that mattered.

In order to test if Moneyball’s basis was correct, I will be using logistic regression to predict whether or not a team will make the playoffs. The seasons that I will analyze go all the way back to 1920, with some missing years in the 70s because of missing values in the contract data. In order to predict if a team makes the playoffs or not, the variables I plan to use are the team’s payroll, the team’s batting average (BA), the team’s earned run average (ERA), their weighted on base average (wOBA) and lastly their fielding independent pitching (FIP). These stats may sound scary but they are fairly simple and will be explained in the data section below.

The data collected is from the **Lahman** package in R (Friendly et al. (2020)). The **Lahman** package contains yearly player statistics going all the way back to the late 1800s. For this project, I will only use the data from 1920 and onward because those years have the most observations for stats like payroll data. In the **Lahman** package there are many different datasets present, for this project I will be using the **Teams** dataset. In order to collect the payroll data, I created a scraper that collected yearly Team stats from Baseball Reference (LLC (n.d.)). The scraper takes the data from each year and saved it into a larger dataset.² The payroll data will be crucial for my report because the basis of Moneyball is fielding the best team by using the least amount of money.

This data analysis was conducted using R (R Core Team (2019)), and in particular the packages Tidyverse (Wickham et al. (2019)) and was compiled using R markdown (Xie, Allaire, and Golemund (2018)). I have 4 sections not including the introduction. The first section covers the data I used for my report, including a few plots to show how the distributions and relationships between various stats. The next section is about the model I used, this will include a in-depth breakdown of logistic regression. Next, I will display the results of my model, this will include the coefficients of the model and an interpretation of the model, along with an application of the model to the results of the 2020 season to see how accurate its predictions are. Lastly, there is a section discussing what we learn from this model and some next steps. Was Moneyball right? Does money impact a team’s ability to make the playoffs? Keep reading to find out!

Data

Model

term	estimate	std.error	statistic	p.value
(Intercept)	-29.5168	2.3831	-12.3858	0.0000
wOBA	137.6176	9.8297	14.0002	0.0000
FIP	-0.2411	0.3711	-0.6496	0.5159
log(payroll)	0.4337	0.0499	8.6986	0.0000
BA	-29.5163	10.7073	-2.7567	0.0058
ERA	-4.9775	0.3524	-14.1248	0.0000

¹Moneyball was actually part of the reason I decided to get into statistics. Applying baseball with numbers seemed like an absolute win to me

²This scraper can be found in the scripts folder of my GitHub repo.




Figure 1: The increase in Payroll over the Years




Figure 2: Relationship between wOBA and Wins




Figure 3: Relationship between FIP and Wins

Table @ref(tab:model) shows us the coefficients from our logistic regression model. We see that teams with a higher wOBA have better chances of making the playoffs, along with team with higher payrolls. This makes sense because when teams will attempt to spend the most money to get the best players and when they score more runs they'll win more games. We can also see that the coefficient for FIP is negative, this does not mean that FIP brings down the chances of making the playoffs. The best values for FIP are the ones closer to 0 and as the value begins to increase, it shows that the pitcher is not that good. What the coefficient is telling us is that lower FIP's give higher chances of making the playoffs while higher FIP's bring down the chances of making the playoffs.




Figure @ref(fig:cooks) shows us the observed cook's distances for the model. Cook's distance gives us the difference between the We see that these cook's distances are very low (most are under 0.01), since our distances are fairly low, we can conclude that the model is fitting the data fairly well.

```
## [1] 0.1898855
```

```
## [1] 0.1593511
```

Results

Discussion

Weaknesses and Next Steps

References

Friendly, Michael, Chris Dalzell, Martin Monkman, and Dennis Murphy. 2020. *Lahman: Sean 'Lahman' Baseball Database*. <https://CRAN.R-project.org/package=Lahman>.

LLC, Sports Reference. n.d. *Baseball-Reference.com*.

R Core Team. 2019. *R: A Language and Environment for Statistical Computing*. Vienna, Austria: R Foundation for Statistical Computing. <https://www.R-project.org/>.

Wickham, Hadley, Mara Averick, Jennifer Bryan, Winston Chang, Lucy D'Agostino McGowan, Romain François, Garrett Golemund, et al. 2019. "Welcome to the tidyverse." *Journal of Open Source Software* 4 (43): 1686. <https://doi.org/10.21105/joss.01686>.

Xie, Yihui, J.J. Allaire, and Garrett Golemund. 2018. *R Markdown: The Definitive Guide*. Boca Raton, Florida: Chapman; Hall/CRC. <https://bookdown.org/yihui/rmarkdown>.